Fake News Detection and Analysis System for Online News and Social Media

Yuxuan Zhang, Jiahui Zhao, Yuchen Sun
NUS-ISS, AIS-PLP, Group 1
Student IDs: A0285664N, A0329852U, A0326248B

I. Business Problem

With the rapid growth of social media and online news platforms, information spreads faster than ever before. However, this has also led to the widespread dissemination of **fake news**, which can mislead the public, influence political opinions, and damage the credibility of legitimate media sources.

For organizations such as media companies, government agencies, and online platforms, manually verifying the authenticity of news content is time-consuming and inefficient. Therefore, there is a strong need for automated systems that can analyze textual content and identify potentially false or misleading information.

This project aims to develop a **Fake News Detection and Analysis System** that automatically analyzes news articles and determines whether they are likely to be fake or real. In addition to detecting fake news, the system will provide supporting analytical information such as topic classification and sentiment analysis to help users better understand the characteristics of the news content.

Such a system can support decision-making processes by enabling users to quickly evaluate the credibility of information circulating online.

II. System Objectives

The main objective of this project is to design and implement an NLP-based system that can analyze news articles and assist users in identifying fake news.

The system aims to:

1) Automatically classify news articles as fake or real.
2) Identify the topic or category of the news article.
3) Analyze the sentiment expressed in the news content.
4) Provide an integrated interface that displays the analysis results to users.

By combining these functionalities, the system will provide a more comprehensive understanding of online news content.

III. System Design

The proposed system will consist of several components that process textual data and perform different natural language processing tasks.

The system workflow is as follows:

1) **User Input**: The user inputs a news article or news text into the system.
2) **Text Preprocessing**: The text is cleaned and prepared for analysis using common preprocessing techniques such as tokenization, normalization, and removal of unnecessary characters.
3) **Fake News Classification Module**: A deep learning model based on BERT (Bidirectional Encoder Representations from Transformers) will be fine-tuned to classify the news article as fake or real.
4) **Topic Classification Module**: The system will identify the main topic of the article, such as politics, health, technology, or business.
5) **Sentiment Analysis Module**: The sentiment of the news content will be analyzed to determine whether the article expresses positive, negative, or neutral sentiment.
6) **Result Presentation**: The system will present the analysis results, including prediction probabilities and supporting information, through a simple interface.

The fake news classifier serves as the core prediction module, while topic classification and sentiment analysis provide supporting analytical signals to help users better interpret the characteristics of the news content.

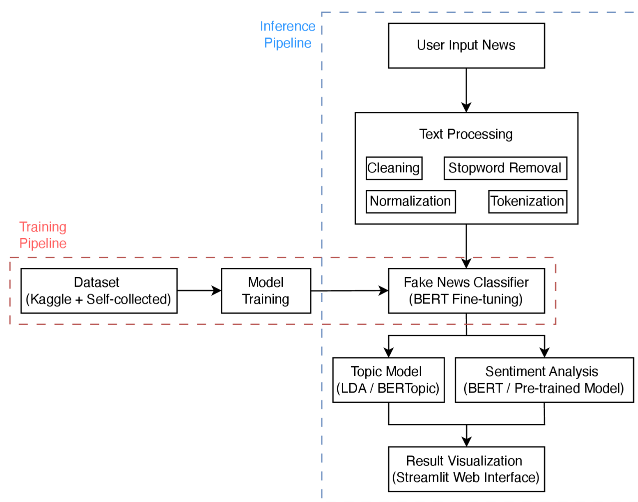

Fig. 1. End-to-end architecture of the proposed fake news detection system, integrating a fine-tuned BERT classifier with topic modeling, sentiment analysis, and a web-based visualization interface.

IV. Dataset Selection

To ensure both training quality and real-world relevance, this project adopts a hybrid dataset strategy that combines a public benchmark with additional self-collected samples.

A. Kaggle Fake News Dataset

The Kaggle Fake News Dataset serves as the primary training source. It provides over 20,000 labeled articles (fake/real), with fields such as title and article text. Its scale and relatively balanced class distribution make it suitable for fine-tuning BERT and establishing a strong baseline.

B. Self-collected News Data

To reduce dataset bias and improve robustness, we will add self-collected articles from reputable news and fact-checking sources (e.g., Snopes and PolitiFact). These samples will be manually labeled or cross-checked to maintain quality and to include more recent writing styles and topics.

We expect to collect approximately 200–500 additional news articles from fact-checking sources to complement the benchmark dataset. For labeling, we will align article claims with fact-check verdicts (e.g., "False", "Pants on Fire" → fake; "True" → real). To improve reliability, at least two team members will cross-validate labels for a subset of samples, and disagreements will be resolved through discussion. We will also deduplicate near-identical articles and ensure that self-collected samples do not overlap with the benchmark test set.

Overall, combining benchmark data with self-collected data provides:

1) sufficient labeled volume for deep learning,
2) diversity for generalization, and
3) more realistic evaluation.

Before training, all data will be standardized through preprocessing (cleaning, tokenization, and normalization).

V. Scope of Work

This section defines what is **included** and **excluded** in the project deliverables.

Included in scope:

- Build a fake news detection module based on a fine-tuned BERT model.
- Implement topic classification and sentiment analysis as supporting analytics.
- Develop an end-to-end pipeline from text input, preprocessing, model inference, to result display.
- Evaluate model performance using accuracy, precision, recall, and F1-score.
- Deliver a simple demo interface for user interaction and result visualization.

Out of scope (for this phase):

- Large-scale production deployment and cloud infrastructure optimization.
- Real-time social media stream monitoring.
- Multimedia fake news detection (images, videos, or audio).

VI. Evaluation

We will split the dataset into training/validation/test sets (e.g., 80/10/10). Model performance will be evaluated primarily using F1-score, along with accuracy, precision, and recall. We will also analyze confusion matrices and conduct qualitative error analysis on misclassified samples to identify common failure cases (e.g., satire, ambiguous claims, or sensational writing).

As a baseline comparison, we may also evaluate a traditional machine learning classifier (e.g., Logistic Regression or SVM with TF-IDF features) to compare with the fine-tuned BERT model.

VII. Effort Estimates

The estimated effort allocation across project stages is as follows.

1) **Dataset preparation (20%)**: Collect, clean, and format datasets for training and testing.
2) **Fake news model development (30%)**: Fine-tune BERT and optimize the core classification pipeline.
3) **Topic and sentiment modules (15%)**: Implement and validate supporting NLP analysis components.
4) **System integration (15%)**: Integrate all modules into one workflow and build a demo interface.
5) **Evaluation and error analysis (10%)**: Measure performance and analyze misclassification patterns.
6) **Documentation and presentation (10%)**: Prepare report, slides, and demonstration materials.

VIII. Risks & Mitigations

To improve project reliability and delivery confidence, we identify key risks and corresponding mitigation strategies.

1) **Data quality and labeling inconsistency**: Self-collected samples may contain noisy labels or ambiguous claims.
   Mitigation: Use clear labeling rules aligned with fact-check verdicts, conduct cross-validation by at least two team members, and resolve conflicts through review meetings.
2) **Class imbalance and domain bias**: The model may overfit to dominant classes or sources and perform poorly on unseen styles.
   Mitigation: Monitor class distribution, apply re-sampling or class weighting when necessary, and evaluate on a held-out, source-diverse test set.
3) **Model overfitting and weak generalization**: Fine-tuned BERT may achieve strong validation scores but degrade on real-world inputs.
   Mitigation: Use early stopping, regularization, and error analysis on misclassified cases (e.g., satire and ambiguous language) to guide iteration.
4) **Integration and schedule slippage**: Combining multiple NLP modules may introduce interface mismatches and delays.
   Mitigation: Define module input/output schemas early, run incremental integration tests, and keep a buffer in the project timeline for debugging.
5) **Interpretability and user trust**: Users may hesitate to rely on predictions without sufficient explanation.
   Mitigation: Present confidence scores and supporting signals (topic and sentiment), and document known limitations clearly in the demo/report.

IX. Conclusion

This proposal presents a practical NLP-based approach to fake news detection with supporting topic and sentiment analysis. By combining a fine-tuned BERT classifier, structured evaluation, and a simple user interface, the project aims to deliver an effective and explainable prototype that helps users assess the credibility of online news content.

The proposed system demonstrates how modern NLP techniques can be applied to assist users in assessing the credibility of online information and understanding the contextual characteristics of potentially misleading news.